

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```


```
In [3]: df = pd.read_csv(r'C:\Users\kriti\OneDrive\Desktop\UM\Bird Strikes Data')
```

```
In [4]: df.head()
```

Out[4]:

	Record ID	Aircraft: Type	Airport: Name	Altitude bin	Aircraft: Make/Model	Wildlife: Number struck	Wildlife: Number Struck Actual	Effect: Impact flight
0	202152	Airplane	LAGUARDIA NY	> 1000 ft	B-737-400	Over 100	859	Engine Struck
1	208159	Airplane	DALLAS/FORT WORTH INTL ARPT	< 1000 ft	MD-80	Over 100	424	No
2	207601	Airplane	LAKEFRONT AIRPORT	< 1000 ft	C-500	Over 100	261	No
3	215953	Airplane	SEATTLE-TACOMA INTL	< 1000 ft	B-737-400	Over 100	806	Precaution: Landing
4	219878	Airplane	NORFOLK INTL	< 1000 ft	CL-RJ100/200	Over 100	942	No

5 rows × 26 columns



```
In [5]: df.shape
```

Out[5]: (25558, 26)

```
In [6]: # check the columns
df.columns
```

```
Out[6]: Index(['Record ID', 'Aircraft: Type', 'Airport: Name', 'Altitude bin',
'Aircraft: Make/Model', 'Wildlife: Number struck',
'Wildlife: Number Struck Actual', 'Effect: Impact to flight',
'FlightDate', 'Effect: Indicated Damage',
'Aircraft: Number of engines?', 'Aircraft: Airline/Operator',
'Origin State', 'When: Phase of flight', 'Conditions: Precipita
tion',
'Remains of wildlife collected?',
'Remains of wildlife sent to Smithsonian', 'Remarks', 'Wildlif
e: Size',
'Conditions: Sky', 'Wildlife: Species',
'Pilot warned of birds or wildlife?', 'Cost: Total $',
'Feet above ground', 'Number of people injured', 'Is Aircraft L
arge?'],
dtype='object')
```

```
In [7]: # information about the dataset
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25558 entries, 0 to 25557
Data columns (total 26 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Record ID                                25558 non-null  int64
1   Aircraft: Type                           25429 non-null  object
2   Airport: Name                            25429 non-null  object
3   Altitude bin                             25429 non-null  object
4   Aircraft: Make/Model                     25558 non-null  object
5   Wildlife: Number struck                  25429 non-null  object
6   Wildlife: Number Struck Actual           25558 non-null  int64
7   Effect: Impact to flight                 25429 non-null  object
8   FlightDate                              25429 non-null  object
9   Effect: Indicated Damage                 25558 non-null  object
10  Aircraft: Number of engines?             25291 non-null  object
11  Aircraft: Airline/Operator               25429 non-null  object
12  Origin State                             25109 non-null  object
13  When: Phase of flight                    25429 non-null  object
14  Conditions: Precipitation                25558 non-null  object
15  Remains of wildlife collected?           25558 non-null  bool
16  Remains of wildlife sent to Smithsonian  25558 non-null  bool
17  Remarks                                  20787 non-null  object
18  Wildlife: Size                           25429 non-null  object
19  Conditions: Sky                          25558 non-null  object
20  Wildlife: Species                        25558 non-null  object
21  Pilot warned of birds or wildlife?       25429 non-null  object
22  Cost: Total $                           25558 non-null  object
23  Feet above ground                       25429 non-null  object
24  Number of people injured                 25558 non-null  int64
25  Is Aircraft Large?                       25429 non-null  object
dtypes: bool(2), int64(3), object(21)
memory usage: 4.7+ MB
```

```
In [8]: # columns with categorical values
df.select_dtypes(include=['object']).columns
```

```
Out[8]: Index(['Aircraft: Type', 'Airport: Name', 'Altitude bin',
              'Aircraft: Make/Model', 'Wildlife: Number struck',
              'Effect: Impact to flight', 'FlightDate', 'Effect: Indicated Da
              mage',
              'Aircraft: Number of engines?', 'Aircraft: Airline/Operator',
              'Origin State', 'When: Phase of flight', 'Conditions: Precipita
              tion',
              'Remarks', 'Wildlife: Size', 'Conditions: Sky', 'Wildlife: Spec
              ies',
              'Pilot warned of birds or wildlife?', 'Cost: Total $',
              'Feet above ground', 'Is Aircraft Large?'],
              dtype='object')
```

```
In [9]: # columns with numerical values
df.select_dtypes(include=['int64', 'float64']).columns
```

```
Out[9]: Index(['Record ID', 'Wildlife: Number Struck Actual',
              'Number of people injured'],
              dtype='object')
```

```
In [10]: # check if there are any null values
df.isnull().values.any() # this function returns true and false
```

```
Out[10]: True
```

```
In [11]: # check how many null values
df.isnull().values.sum()
```

```
Out[11]: 7035
```

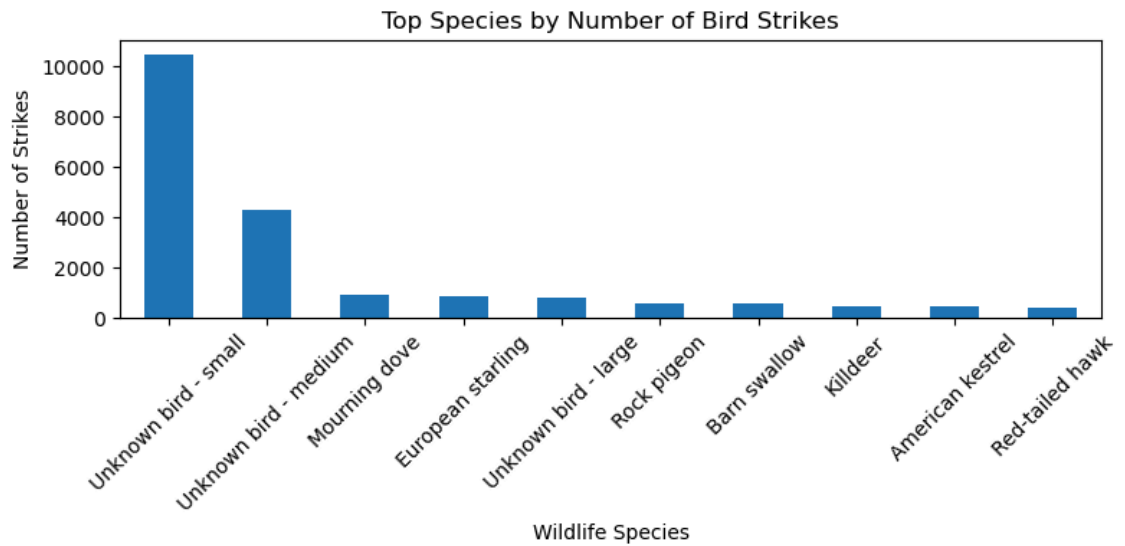
```
In [12]: # Drop rows with null values
df_cleaned = df.dropna()

# Fill null values with a specific value
df_filled = df.fillna(0) # Replace null values with 0
```

```
In [13]: # Count the number of bird strikes by Wildlife Species
bird_strikes_count = df['Wildlife: Species'].value_counts()
```

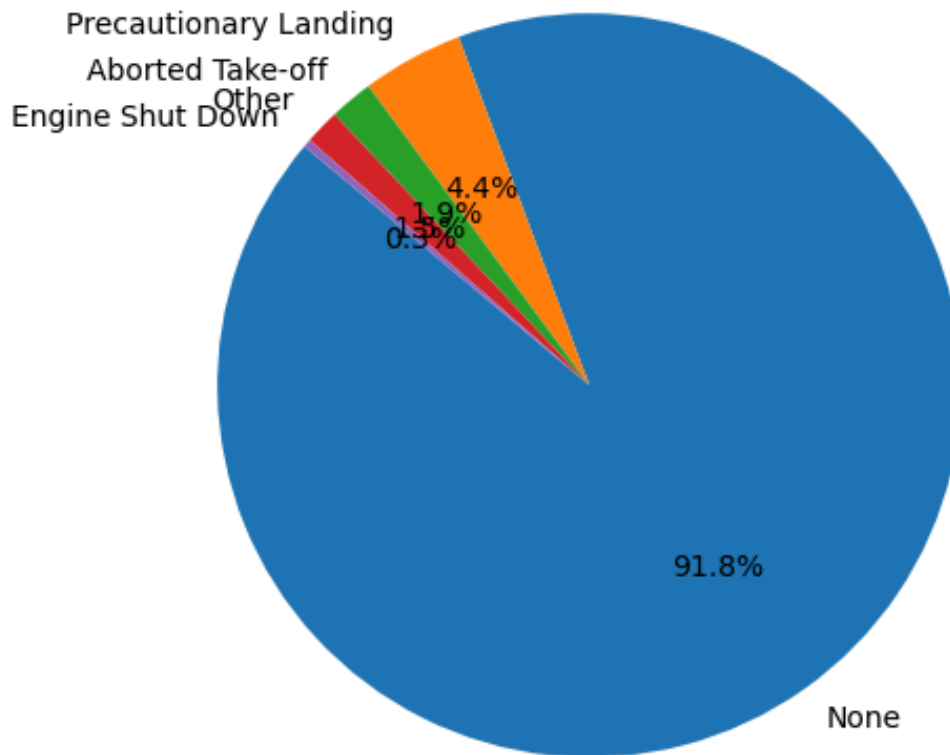
```
In [14]: # Choose the top N species with the highest number of strikes
top_species_count = bird_strikes_count.head(10) # Change 10 to the des
```

```
In [18]: # Plot a bar chart for the top N species
plt.figure(figsize=(8, 4))
top_species_count.plot(kind='bar')
plt.title('Top Species by Number of Bird Strikes')
plt.xlabel('Wildlife Species')
plt.ylabel('Number of Strikes')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



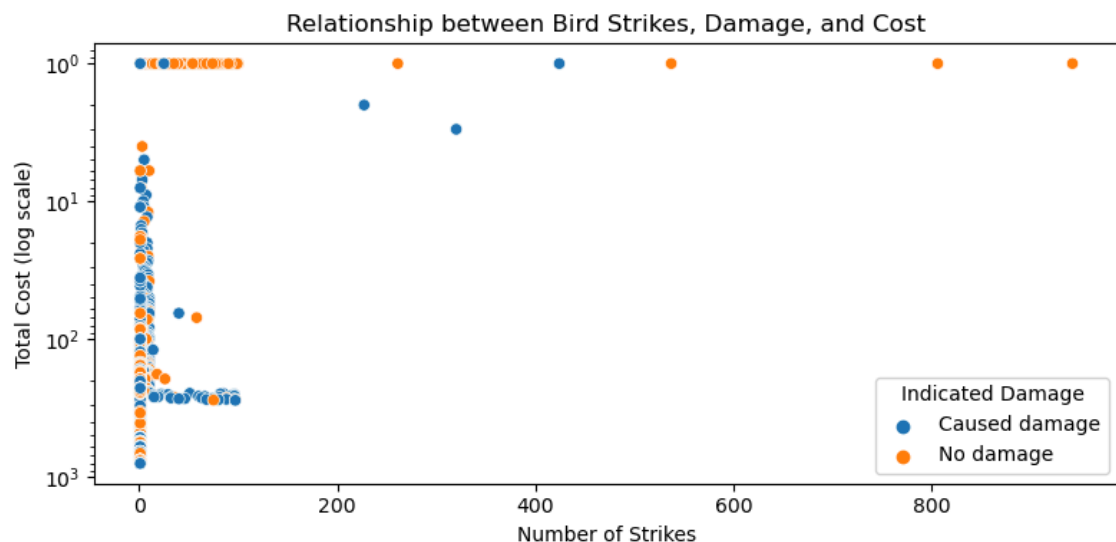
```
In [19]: # Plot a pie chart for the distribution of bird strikes by Effect: Impact
impact_counts = df['Effect: Impact to flight'].value_counts()
plt.figure(figsize=(5, 5))
plt.pie(impact_counts, labels=impact_counts.index, autopct='%1.1f%%', s
plt.title('Distribution of Bird Strikes by Impact to Flight')
plt.axis('equal')
plt.tight_layout()
plt.show()
```

Distribution of Bird Strikes by Impact to Flight



In [21]:

```
# Plot a scatter plot with log scale for the y-axis
plt.figure(figsize=(8, 4))
sns.scatterplot(data=df, x='Wildlife: Number Struck Actual', y='Cost: Total $')
plt.title('Relationship between Bird Strikes, Damage, and Cost')
plt.xlabel('Number of Strikes')
plt.ylabel('Total Cost (log scale)')
plt.yscale('log') # Set y-axis to log scale
plt.legend(title='Indicated Damage')
plt.tight_layout()
plt.show()
```



When using a log scale, the values on the axis are not evenly spaced like in a linear scale. Instead, each increment on the axis represents a multiple of the base of the logarithm. For example, in a logarithmic base 10 scale, each increment on the y-axis represents a multiple of 10.

The main benefit of using a log scale in a scatter plot is that it can help in visualizing data that has a wide range of values or data points that are clustered at low values with a few extreme outliers. It compresses the higher values, making the plot more interpretable and highlighting patterns that may not be apparent on a linear scale.

In the context of the scatter plot code provided earlier, setting the y-axis to a log scale (`plt.yscale('log')`) adjusts the scale of the y-axis to better accommodate the wide range of values in the 'Cost: Total \$' column, making it easier to read and interpret the data points.

```
In [18]: # Filter data for US airlines
us_airlines = df[df['Origin State'] == 'US']
```


```
In [35]: # Top 10 US Airlines with most bird strikes
top_10_airlines = us_airlines.groupby('Aircraft: Airline/Operator')['Wi
```

```
In [36]: # Convert 'Wildlife: Number struck' column to numeric
df['Wildlife: Number struck'] = pd.to_numeric(df['Wildlife: Number struck'])

# Filter data for incidents with non-null 'Wildlife: Number struck'
df_filtered = df.dropna(subset=['Wildlife: Number struck'])

# Group by 'Airport: Name' and sum the 'Wildlife: Number struck' for each airport
top_airports = df_filtered.groupby('Airport: Name')['Wildlife: Number struck'].sum()

# Plotting Top 50 Airports with most bird strikes
plt.figure(figsize=(12, 8))
top_airports.plot(kind='barh')
plt.title('Top 50 Airports with most Bird Strikes')
plt.xlabel('Number of Bird Strikes')
plt.ylabel('Airport Name')
plt.show()
```




```
In [38]: # Convert 'FlightDate' column to datetime
df['FlightDate'] = pd.to_datetime(df['FlightDate'], errors='coerce')

# Filter data for incidents with non-null 'FlightDate'
df_filtered = df.dropna(subset=['FlightDate'])

# Extract month from 'FlightDate' and create a new 'Month' column
df_filtered['Month'] = df_filtered['FlightDate'].dt.month

# Group by month and count the number of incidents
monthly_incidents = df_filtered.groupby('Month').size()

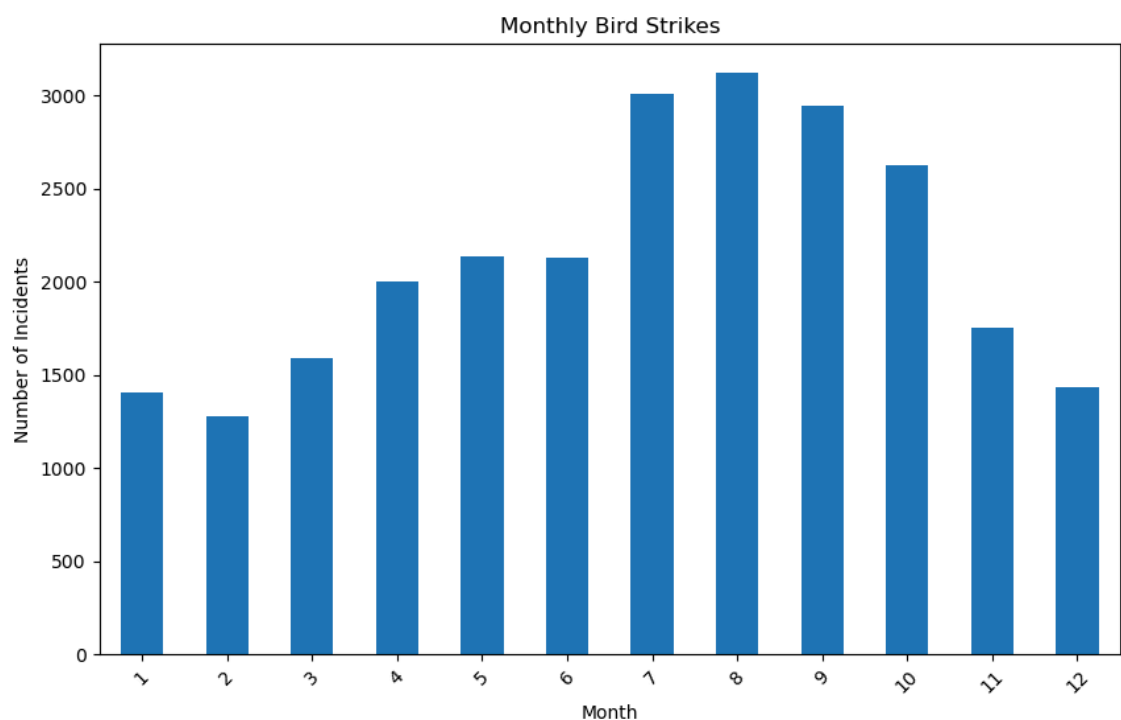
# Plotting Monthly Bird Strikes
plt.figure(figsize=(10, 6))
monthly_incidents.plot(kind='bar')
plt.title('Monthly Bird Strikes')
plt.xlabel('Month')
plt.ylabel('Number of Incidents')
plt.xticks(rotation=45)
plt.show()
```

C:\Users\kriti\AppData\Local\Temp\ipykernel_18176\1838022889.py:8: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

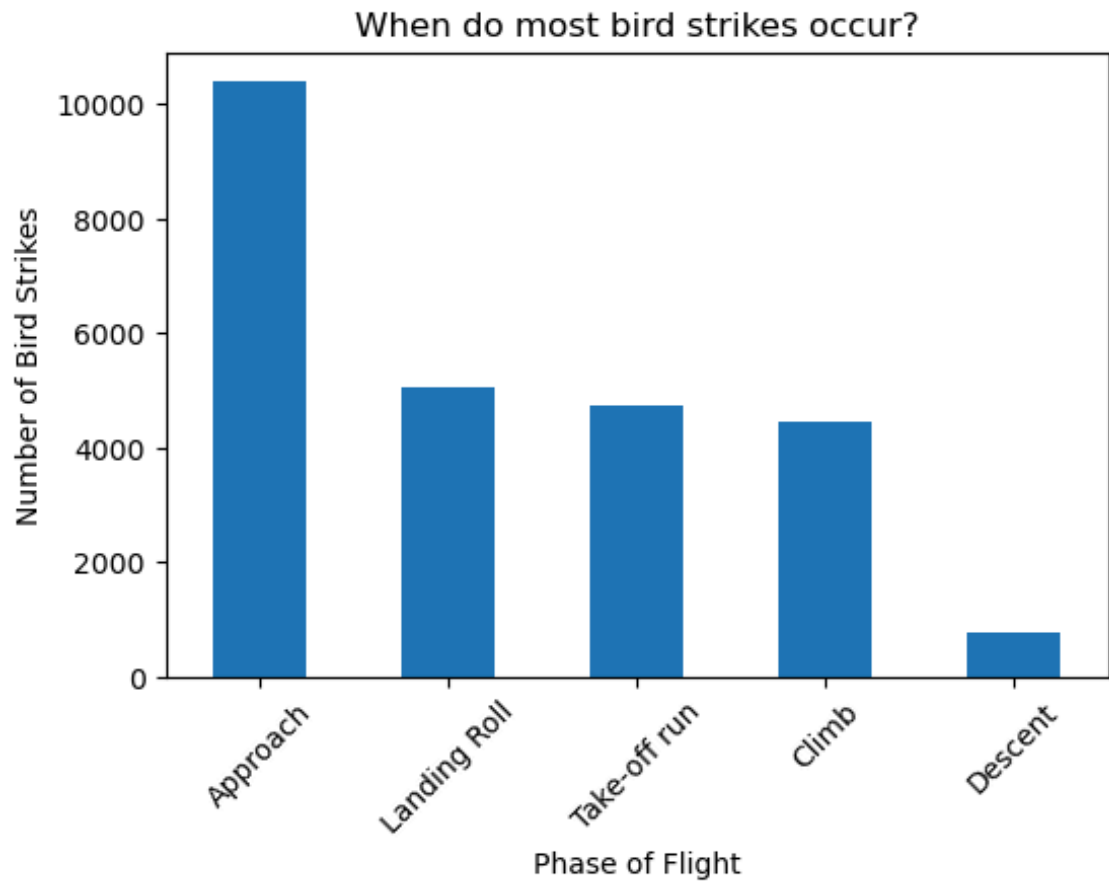
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
(https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_filtered['Month'] = df_filtered['FlightDate'].dt.month
```



```
In [23]: # When do most bird strikes occur?
most_strikes_phase = df['When: Phase of flight'].value_counts().nlargest(5)

# Plotting When do most bird strikes occur
plt.figure(figsize=(6, ))
most_strikes_phase.plot(kind='bar')
plt.title('When do most bird strikes occur?')
plt.xlabel('Phase of Flight')
plt.ylabel('Number of Bird Strikes')
plt.xticks(rotation=45)
plt.show()
```



```
In [25]: # Altitude of airplanes at the time of strike
altitude_strike = df.groupby('Altitude bin')['Wildlife: Number struck']

# Plotting Altitude of airplanes at the time of strike
plt.figure(figsize=(6, 6))
altitude_strike.plot(kind='bar')
plt.title('Altitude of airplanes at the time of Bird Strikes')
plt.xlabel('Altitude Bin')
plt.ylabel('Number of Bird Strikes')
plt.xticks(rotation=45)
plt.show()
```


TypeError

Traceback (most recent call

last)

Cell In[25], line 6

```
4 # Plotting Altitude of airplanes at the time of strike
5 plt.figure(figsize=(6, 6))
----> 6 altitude_strike.plot(kind='bar')
7 plt.title('Altitude of airplanes at the time of Bird Strikes')
8 plt.xlabel('Altitude Bin')
```

File E:\Anaconda\Lib\site-packages\pandas\plotting_core.py:1000, in PlotAccessor.__call__(self, *args, **kwargs)

```
997         label_name = label_kw or data.columns
998         data.columns = label_name
-> 1000 return plot_backend.plot(data, kind=kind, **kwargs)
```

File E:\Anaconda\Lib\site-packages\pandas\plotting_matplotlib__init__.py:71, in plot(data, kind, **kwargs)

```
69         kwargs["ax"] = getattr(ax, "left_ax", ax)
70 plot_obj = PLOT_CLASSES[kind](data, **kwargs)
----> 71 plot_obj.generate()
72 plot_obj.draw()
73 return plot_obj.result
```


File E:\Anaconda\Lib\site-packages\pandas\plotting_matplotlib\core.py:450, in MPLPlot.generate(self)

```
448 def generate(self) -> None:
449     self._args_adjust()
--> 450     self._compute_plot_data()
451     self._setup_subplots()
452     self._make_plot()
```

File E:\Anaconda\Lib\site-packages\pandas\plotting_matplotlib\core.py:635, in MPLPlot._compute_plot_data(self)

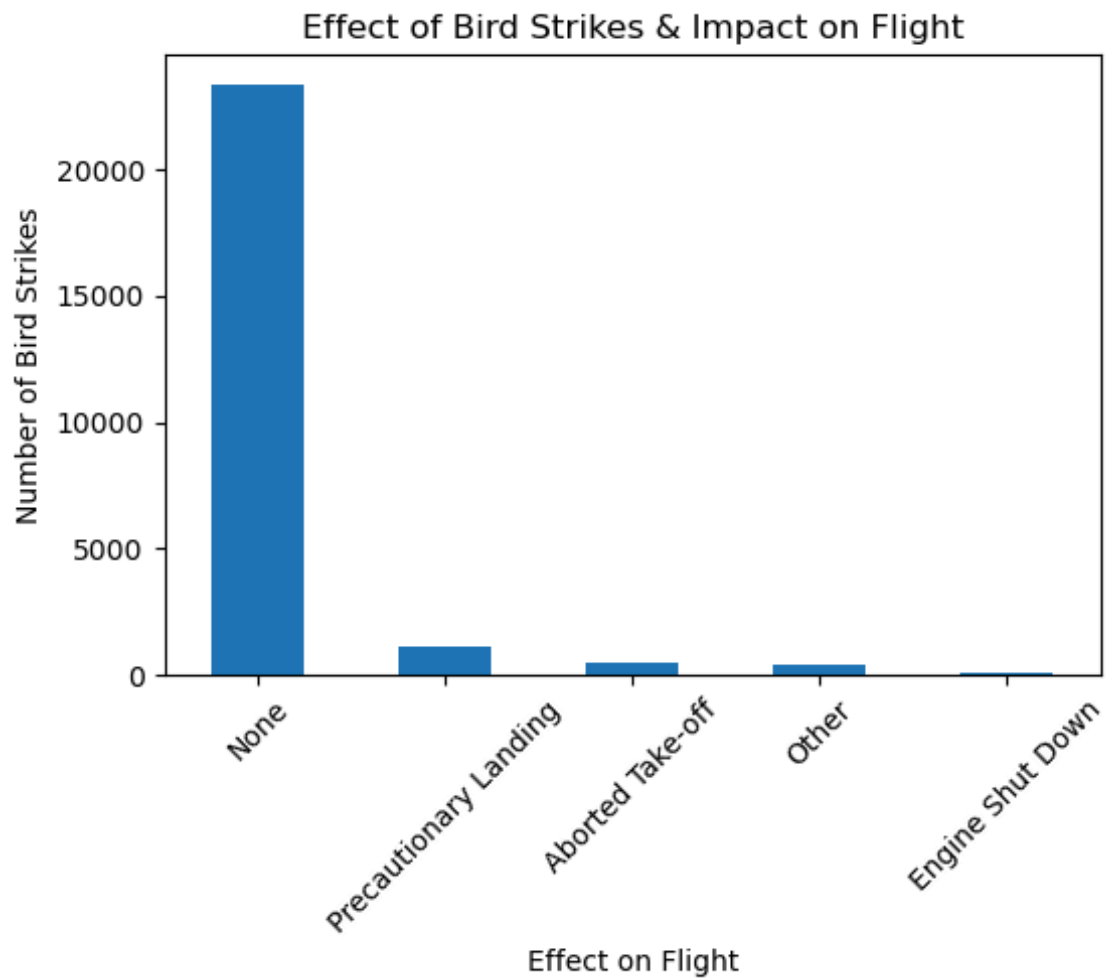
```
633 # no non-numeric frames or series allowed
634 if is_empty:
--> 635     raise TypeError("no numeric data to plot")
637 self.data = numeric_data.apply(self._convert_to_ndarray)
```

TypeError: no numeric data to plot



```
In [26]: # Effect of Bird Strikes & Impact on Flight
effect_impact = df['Effect: Impact to flight'].value_counts()


# Plotting Effect of Bird Strikes & Impact on Flight
plt.figure(figsize=(6, 4))
effect_impact.plot(kind='bar')
plt.title('Effect of Bird Strikes & Impact on Flight')
plt.xlabel('Effect on Flight')
plt.ylabel('Number of Bird Strikes')
plt.xticks(rotation=45)
plt.show()
```



```
In [42]: # Were Pilots Informed? & Prior Warning and Effect of Strike Relation
pilot_warning = df['Pilot warned of birds or wildlife?'].value_counts()

# Plotting Were Pilots Informed? & Prior Warning and Effect of Strike R
plt.figure(figsize=(6, 4))
pilot_warning.plot(kind='pie', autopct='%1.1f%%')
plt.title('Were Pilots Informed? & Prior Warning and Effect of Strike R
plt.ylabel('')
plt.show()
```

Were Pilots Informed? & Prior Warning and Effect of Strike Relation



```
In [27]: # Import necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report

# Data preprocessing
# Encode categorical variables
le = LabelEncoder()
df['Effect: Impact to flight'] = le.fit_transform(df['Effect: Impact to
df['When: Phase of flight'] = le.fit_transform(df['When: Phase of flight'])

# Select features and target variable
X = df[['Effect: Impact to flight', 'When: Phase of flight']]
y = df['Wildlife: Number Struck Actual']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

# Initialize the Random Forest Classifier
rf_classifier = RandomForestClassifier()

# Train the classifier
rf_classifier.fit(X_train, y_train)

# Make predictions
y_pred = rf_classifier.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Generate classification report
print(classification_report(y_test, y_pred))
```

Accuracy: 0.8151408450704225

	precision	recall	f1-score	support
1	0.82	1.00	0.90	4167
2	0.00	0.00	0.00	112
3	0.00	0.00	0.00	99
4	0.00	0.00	0.00	114
5	0.00	0.00	0.00	90
6	0.00	0.00	0.00	104
7	0.00	0.00	0.00	92
8	0.00	0.00	0.00	90
9	0.00	0.00	0.00	96
10	0.00	0.00	0.00	87
22	0.00	0.00	0.00	1
28	0.00	0.00	0.00	1
32	0.00	0.00	0.00	1
33	0.00	0.00	0.00	1
35	0.00	0.00	0.00	4
36	0.00	0.00	0.00	1
39	0.00	0.00	0.00	4
44	0.00	0.00	0.00	2
46	0.00	0.00	0.00	1
51	0.00	0.00	0.00	1
52	0.00	0.00	0.00	2
56	0.00	0.00	0.00	2
58	0.00	0.00	0.00	1
59	0.00	0.00	0.00	1
60	0.00	0.00	0.00	3
61	0.00	0.00	0.00	1
62	0.00	0.00	0.00	3
63	0.00	0.00	0.00	1
64	0.00	0.00	0.00	1
66	0.00	0.00	0.00	1
68	0.00	0.00	0.00	1
69	0.00	0.00	0.00	1
72	0.00	0.00	0.00	1
73	0.00	0.00	0.00	5
74	0.00	0.00	0.00	2
75	0.00	0.00	0.00	1
76	0.00	0.00	0.00	1
77	0.00	0.00	0.00	2
78	0.00	0.00	0.00	1
79	0.00	0.00	0.00	1
80	0.00	0.00	0.00	1
81	0.00	0.00	0.00	1
82	0.00	0.00	0.00	1
84	0.00	0.00	0.00	1
85	0.00	0.00	0.00	1
88	0.00	0.00	0.00	1
89	0.00	0.00	0.00	1
92	0.00	0.00	0.00	1
95	0.00	0.00	0.00	1
96	0.00	0.00	0.00	1
98	0.00	0.00	0.00	1
99	0.00	0.00	0.00	1
accuracy			0.82	5112
macro avg	0.02	0.02	0.02	5112

weighted avg	0.66	0.82	0.73	5112
--------------	------	------	------	------

```
E:\Anaconda\Lib\site-packages\sklearn\metrics\_classification.py:1469:
UndefinedMetricWarning: Precision and F-score are ill-defined and bein
g set to 0.0 in labels with no predicted samples. Use `zero_division`
parameter to control this behavior.
```

```
    _warn_prf(average, modifier, msg_start, len(result))
```

```
E:\Anaconda\Lib\site-packages\sklearn\metrics\_classification.py:1469:
UndefinedMetricWarning: Precision and F-score are ill-defined and bein
g set to 0.0 in labels with no predicted samples. Use `zero_division`
parameter to control this behavior.
```

```
    _warn_prf(average, modifier, msg_start, len(result))
```

```
E:\Anaconda\Lib\site-packages\sklearn\metrics\_classification.py:1469:
UndefinedMetricWarning: Precision and F-score are ill-defined and bein
g set to 0.0 in labels with no predicted samples. Use `zero_division`
parameter to control this behavior.
```

```
    _warn_prf(average, modifier, msg_start, len(result))
```

In []:

In []:

In []: